Claude Cowork and Agentic Workflows: AI That Works With You

Learning Objectives

By the end of this session, you should be able to:

- Define Claude Cowork and agentic workflows, and recall how they differ from chat-based prompting.
- Explain what changes when moving from Chat to Cowork (task definition vs. manual stitching; end-to-end execution).
- Describe key Cowork capabilities (local file access, sub-agent coordination, professional outputs, scheduled tasks, context continuity, plugins).
- Execute a workflow using the Agentic Loop: Plan → Approve → Execute → Iterate.
- Identify where human oversight must be inserted (visible reasoning, approval gates, supervisor role) to ensure safe professional use.

Learning Objectives

By the end of this session, you should be able to:

- Create and reuse Claude Skills to standardize repeatable workflows and improve output consistency across a team.
- Assess an agentic workflow's effectiveness using evidence from outputs (quality, repeatability, risk controls, time saved).

Agenda

In this session, we'll discuss:

- Transition From Chat to Cowork
- Key Cowork Capabilities
- Agentic Loop Framework: Planning, Approval, Execution, and Iteration
- Human Oversight Mechanisms: Visible Reasoning, Approval Gate, and Supervisory Roles
- Claude Skill Set
- Cowork in Action: Monthly Close Pack Case Study
- Agentic Characteristics

From Chat to Cowork

What Actually Changes When You Move Beyond the Chat Interface

What Actually Changes

In Chat

You prompt, Claude responds, and you stitch outputs together manually.

In Cowork

You define a task, and Claude executes it as an agent, end-to-end.

This is not a better chatbot—it's a desktop AI agent for knowledge work

Key Cowork Capabilities

Direct Local File Access	• Reads from and writes to local files—no manual uploads or downloads.
Sub-Agent Coordination	• For complex tasks, breaks work into smaller tasks and coordinates parallel workstreams.
Professional Outputs	• Generates Excel spreadsheets, PowerPoint presentations, and formatted documents.
Scheduled Tasks	• Type /schedule to run tasks on demand or on a cadence; runs while Claude Desktop is open.
Context Continuity	• Remembers the task, intermediate outputs, and direction across all steps.
Plugins	• Bundle skills, connectors, and sub-agents into a single package for role, team, or company.

The Agentic Loop

Plan → Approve → Execute → Iterate

Plan → Approve → Execute → Iterate

1 Plan

- Claude reads context and announces what it will do.

2 Approve

- Human reviews the plan before Claude starts.

3 Execute

- Claude works: reading files, analyzing, and generating.

4 Iterate

- Review output, suggest refinements, and Claude adjusts.

↻ **Like delegating to a competent colleague—brief them, approve the approach, let them work, and review the result.**

Human Oversight

The Guardrails That Make Agentic AI Safe for Professional Work

Human Oversight in Cowork

Visible Reasoning

Claude shows its plan and thinking before and during tasks; you always know what it's doing and why.

Approval Gates

Claude asks before taking significant actions like creating files or modifying data.

You Are the Supervisor

Claude proposes, you approve.

Nothing happens without your sign-off.

This is what makes agentic AI safe for real professional work.

Claude Skills

From One-Off Workflows to Reusable, Team-Wide Automation

Reusable Workflows With Skills

What Are Skills?

Deterministic, reusable workflow instructions that can be triggered repeatedly.

Once you've created a workflow in Cowork, save it as a Skill for repeated use.

Instead of reusing prompts (Chat), you reuse entire workflows (Cowork).

Benefits

Repeatability: Run the same workflow consistently.

Output Consistency: Same quality every time.

Team-Wide Usage: Share skills across your team.

Progression: CLAUDE.md (project context) → Skills (reusable across projects and teams).

Cowork in Action

The Monthly Close Pack—From Five Source Files to a Finished Deliverable

The Scenario: Monthly Close Pack

Monday, April 7, 2025, 8:15 a.m. ExCo meeting at 9 a.m.
The Operations Director needs the monthly close pack on the table. It pulls from five separate data sources across five folders. Normally: 2.5–3 hours of analyst work.

Five Source Files (Local, No Uploads)	CLAUDE.md—Persistent Context
• finance/Q1_Actuals_March_Final.xlsx • sales/CRM_Pipeline_March_2025.csv • support/Zendesk_Tickets_March_2025.csv • product/Product_Velocity_March_2025.csv • hr/Headcount_March_2025.csv	• Folder structure and file naming convention • KPI targets for all five functions • Flag thresholds (e.g., CSAT < 80% = RED) • Eight-step workflow Claude must follow • Standing rules: citations, number format • Output folder convention **Loads automatically every session—Claude already knows the rules.**

The Prompt and What Cowork Finds

The entire user prompt: "Run the April monthly close pack using the March data files in this project. Save everything to outputs/Monthly_Close_April2025/"

Three red flags Claude identifies independently (from the source data):

Pipeline Coverage Ratio

- Target: ≥ 3.0× | Actual: 2.28× | Gap: (0.72×)
- Two Stage 4 deals stalling past 30 days. VP Sales to provide plan by 14 April.

Blended CSAT Score

- Target: ≥ 85.0% | Actual: 81.6% | Gap: (3.4pp)
- Product/Bug (72.5%) and Latency (70.0%) dragging average. P1 post-mortem by 10 April.

PS Delivery Headcount

- Target: ≥ 10 FTE | Actual: 7 FTE | Gap: (3 FTE)
- Root cause of £375k Q1 cost overrun. 2 JDs posted, no offers.

Cowork's Eight Agentic Steps

Triggered by one short prompt. Claude narrates every step and waits for approval before starting.

1 Load CLAUDE.md and announce plan	**2** Read all 5 files in parallel	**3** Cross-source anomaly detection	**4** Build KPI Dashboard Excel
5 Write Management Commentary (Word)	**6** Write Red Flags Memo	**7** Create output folder and save	**8** /schedule: make it recurring

Outputs: KPI Dashboard (Excel) | Management Commentary (Word) | Red Flags Memo → All saved to outputs/

What Makes This Agentic?

Each behaviour below is something Claude does autonomously that a reactive chat assistant cannot do.

No File Upload	• Claude opens local files directly from disk. The entire five-file read happens without the user touching a single file.
CLAUDE.md Auto-Load	• Targets, thresholds, folder structure, and standing rules load automatically. Claude already knows what "good" looks like.
Plan-Before-Act	• Claude announces the full eight-step plan before writing a single file. The user can inspect, redirect, or stop.
Parallel Sub-Agents	• All five source files are read as parallel workstreams. This is sub-agent coordination, not sequential processing.
Cross-Source Synthesis	• The three red flags each come from a different file. The commentary connects the headcount CSV to the Excel cost overrun.
/schedule	• One command turns a one-off task into a recurring autonomous workflow. Next month's close pack runs without user input.

Summary

Here's a brief recap:

- Cowork is not a better chatbot it acts like a desktop AI agent for knowledge work, where you define a task and Claude executes it end-to-end.
- Cowork enables multi-step work through: direct local file access, parallel sub-agent coordination, professional deliverables (Excel/PowerPoint/Word), scheduled runs, and context continuity across steps.
- The Agentic Loop mirrors delegating to a competent colleague: Claude proposes a plan, the human approves, Claude executes, and the work is refined through feedback.
- Human oversight is the safety layer: Claude provides visible reasoning, pauses at approval gates, and keeps the user as supervisor nothing significant happens without sign-off.

Summary

Here's a brief recap:

- Claude Skills turn one-off workflows into reusable, deterministic playbooks, improving repeatability, output consistency, and team-wide reuse.
- In the Monthly Close Pack example, Cowork reads multiple source files, detects cross-source anomalies, produces deliverables, saves outputs to folders, and can schedule the workflow to recur.

Learning Outcomes

You should now be able to:

- Identify core terms (Cowork, agentic workflow, Skills, approval gates) and how Cowork differs from Chat.
- Summarize how Cowork enables multi-step execution using files, persistent context, and automation features.
- Explain the value and risks of agentic workflows and the role of Human-in-the-Loop controls.
- Run a multi-step workflow using Plan → Approve → Execute → Iterate, including providing approvals and feedback at the right moments.
- Decide where oversight is required based on task risk and impact.

Learning Outcomes

You should now be able to:

- Build and share a skill to operationalize a repeatable workflow for a team (consistent structure, outputs, and rules).
- Review workflow outputs and controls to judge whether the workflow is reliable, repeatable, and safe for professional use.